

Set up

```
# Du code pour vérifier la VM sur laquelle je tourne
from psutil import virtual_memory
ram_gb = virtual_memory().total / 1e9
print('Your runtime has {:.1f} gigabytes of available RAM\n'.format(ram_gb))
```

```
if ram_gb < 20:
    print('Not using a high-RAM runtime')
else:
    print('You are using a high-RAM runtime!')
```

↗ Your runtime has 359.2 gigabytes of available RAM

You are using a high-RAM runtime!

```
import os
import pandas as pd
import numpy as np
import math
import matplotlib.pyplot as plt
from google.colab import drive
```

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score, recall_score, precision_score
from sklearn.linear_model import LogisticRegression
```

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
```

```
drive.mount('/content/drive')
```

↗ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True)

Loading dataset

Ce dataset provient de : <https://data.mendeley.com/datasets/c4n7fckkz3/3> Avec l'article associé :

<https://link.springer.com/article/10.1007/s10207-023-00723-w>

```
folder_path="/content/drive/MyDrive/DNS_Exfiltration_Dataset /DNS_Exfiltration_Dataset/"
csv_file_path = os.path.join(folder_path, 'dataset.csv')
data = pd.read_csv(csv_file_path)
```

```
# Load the English words from file
with open(os.path.join(folder_path, 'english_words.txt'), 'r') as f:
    english_words = set(f.read().splitlines())
```

```
print(english_words)
```

↗ curriculum', 'consecrates', 'firs', 'usurp', 'downscaling', 'playtime', 'robbery', 'spottiest', 'subverts', 'hullos', 'fr

```
# Adjust display options to ensure full content is shown
```

```
pd.set_option('display.max_rows', 100)
pd.set_option('display.max_columns', None)
pd.set_option('display.max_colwidth', None)
```

```
data.info()
```

↗

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 35074150 entries, 0 to 35074149
Data columns (total 22 columns):
 #   Column                Dtype
---  -
 0   186.169.253.58         object
 1   surbl.org              object
 2   1624438272607          int64
 3   False                 bool
 4   h.surbl.org            object
 5   1                     int64
 6   1.1                   int64
```

```

7  0          int64
8  0.1        int64
9  -0.0       float64
10 0.0        float64
11 0.0.1      float64
12 0.0.2      float64
13 0.0.3      float64
14 3.444444444444446 float64
15 9.59311095410544 float64
16 1.5        float64
17 1.5811388300841898 float64
18 468.75     float64
19 0.4444444444444444 float64
20 0.25849625007211563 float64
21 0.81743691684035 float64
dtypes: bool(1), float64(13), int64(5), object(3)
memory usage: 5.5+ GB

# On nomme les colonnes (d'après la documentation)
data.columns = [
    'user_ip', 'domain', 'timestamp', 'attack', 'request', 'len',
    'subdomains_count', 'w_count', 'w_max', 'entropy', 'w_max_ratio',
    'w_count_ratio', 'digits_ratio', 'uppercase_ratio', 'time_avg',
    'time_stdev', 'size_avg', 'size_stdev', 'throughput', 'unique',
    'entropy_avg', 'entropy_stdev'
]

```

```

# Vérification
print(data.columns)
data.info()

```

```

Index(['user_ip', 'domain', 'timestamp', 'attack', 'request', 'len',
      'subdomains_count', 'w_count', 'w_max', 'entropy', 'w_max_ratio',
      'w_count_ratio', 'digits_ratio', 'uppercase_ratio', 'time_avg',
      'time_stdev', 'size_avg', 'size_stdev', 'throughput', 'unique',
      'entropy_avg', 'entropy_stdev'],
      dtype='object')
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 35074150 entries, 0 to 35074149
Data columns (total 22 columns):
#   Column                Dtype
---  -
0   user_ip               object
1   domain                object
2   timestamp             int64
3   attack                bool
4   request               object
5   len                   int64
6   subdomains_count      int64
7   w_count                int64
8   w_max                 int64
9   entropy                float64
10  w_max_ratio            float64
11  w_count_ratio          float64
12  digits_ratio           float64
13  uppercase_ratio        float64
14  time_avg               float64
15  time_stdev             float64
16  size_avg               float64
17  size_stdev             float64
18  throughput             float64
19  unique                 float64
20  entropy_avg            float64
21  entropy_stdev          float64
dtypes: bool(1), float64(13), int64(5), object(3)
memory usage: 5.5+ GB

```

```

# Check the first few rows of the dataset
print(data.head())

```

```

user_ip      domain      timestamp  attack \
0  186.169.253.58  surbl.org  1624438272607  False
1  186.169.253.58  spamhaus.org  1624438273058  False
2  186.169.253.58  spamhaus.org  1624438273058  False
3  186.169.253.58  spamhaus.org  1624438273059  False
4  186.169.253.58  spamhaus.org  1624438273059  False

request  len  subdomains_count  w_count  w_max \
0  f.surbl.org  1  1  0  0
1  118.141.11.106.sbl.spamhaus.org  18  5  0  0
2  118.141.11.106.zen.spamhaus.org  18  5  1  3
3  128.141.11.106.sbl.spamhaus.org  18  5  0  0
4  128.141.11.106.zen.spamhaus.org  18  5  1  3

entropy  w_max_ratio  w_count_ratio  digits_ratio  uppercase_ratio \
0 -0.000000  0.000000  0.000000  0.000000  0.0

```

1	2.633731	0.000000	0.000000	0.611111	0.0
2	2.633731	0.166667	0.055556	0.611111	0.0
3	2.863826	0.000000	0.000000	0.611111	0.0
4	2.863826	0.166667	0.055556	0.611111	0.0

	time_avg	time_stdev	size_avg	size_stdev	throughput	unique \
0	0.222222	0.440959	1.0	0.000000	3333.333333	0.555556
1	55.555556	165.542375	17.2	0.421637	343.313373	0.000000
2	0.333333	0.500000	17.2	0.421637	43000.000000	0.000000
3	0.333333	0.500000	17.3	0.483046	43250.000000	0.000000
4	0.333333	0.500000	17.4	0.516398	43500.000000	0.000000

	entropy_avg	entropy_stdev
0	0.000000	0.000000
1	3.048277	0.177285
2	2.983547	0.199622
3	2.959741	0.198131
4	2.935936	0.193400

```
# Check for any missing values
print(data.isnull().sum())
```

```
user_ip      0
domain       0
timestamp    0
attack       0
request      0
len          0
subdomains_count 0
w_count      0
w_max        0
entropy      0
w_max_ratio  0
w_count_ratio 0
digits_ratio 0
uppercase_ratio 0
time_avg     0
time_stdev   0
size_avg     0
size_stdev   0
throughput   0
unique       0
entropy_avg  0
entropy_stdev 0
dtype: int64
```

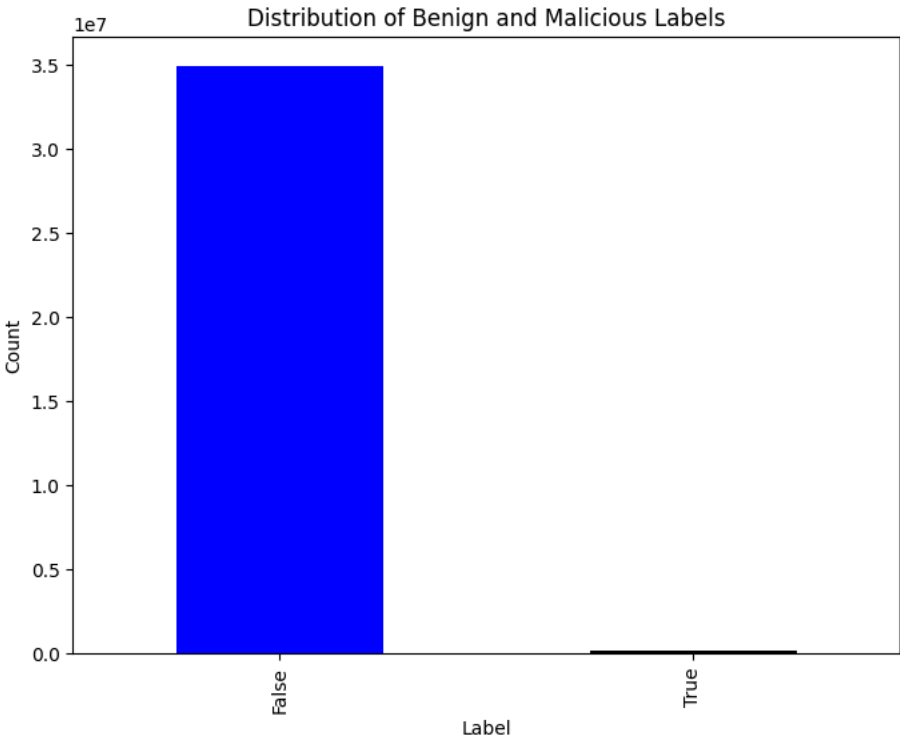
Visualisation du dataset

```
# Count the number of attacks
print(data['attack'].value_counts())
```

```
num_benign, num_attacks = data['attack'].value_counts()
ratio = num_attacks / (num_attacks + num_benign)
print("ratio of attacks in the dataset: ", ratio*100,"%")
```

```
attack
False    34899371
True      174779
Name: count, dtype: int64
ratio of attacks in the dataset:  0.49831286004079933 %
```

```
plt.figure(figsize=(8, 6))
data["attack"].value_counts().plot(kind='bar', color=['blue', 'black'])
plt.title('Distribution of Benign and Malicious Labels')
plt.xlabel('Label')
plt.ylabel('Count')
plt.show()
```

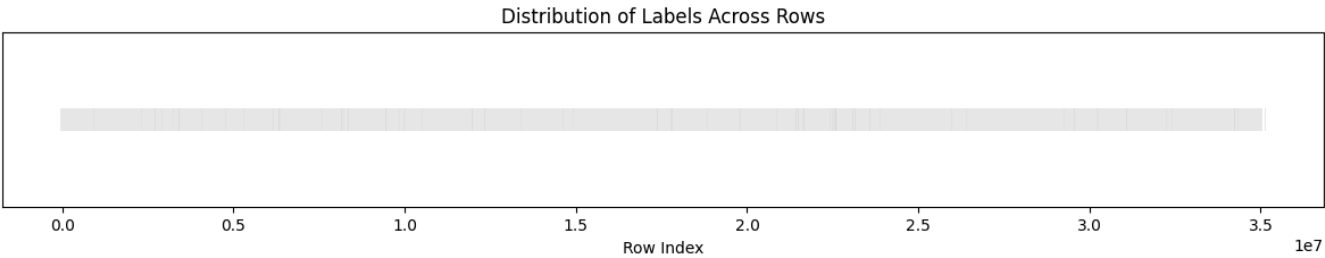


```
# Create a color map: True will be green, False will be red
color_map = data['attack'].map({True: 'black', False: 'white'})

# Plotting
plt.figure(figsize=(15, 2))

# Scatter plot where each point's color represents the label
plt.scatter(data.index, [1]*len(data), c=color_map, marker='|', s=200)

plt.xlabel('Row Index')
plt.yticks([])
plt.title('Distribution of Labels Across Rows')
plt.show()
```



```
#Printing some random rows
random_rows = data.sample(n=10, random_state=42)
print(random_rows)
```



	user_ip	domain	timestamp	attack	\
24577790	186.169.157.176	spamhaus.org	1624517077166	False	
33959090	130.9.48.128	root-servers.net	1624531070154	False	
29079969	130.9.40.66	local.local	1624523951775	False	
32988913	186.169.134.52	bg.ac.rs	1624529568649	False	
18346957	130.9.40.66	local.local	1624491628548	False	
26706171	130.9.48.179	googlesyndication.com	1624520513843	False	
15605255	6c47:8c4e:2567::62e7	bg.ac.rs	1624479125869	False	
27688005	130.9.47.63	root-servers.net	1624521973857	False	
24907918	130.9.40.66	local.local	1624517631226	False	
17026505	130.9.40.66	local.local	1624485232008	False	

	request	len	subdomains_count	w_count	\
24577790	7.22.142.174.zen.spamhaus.org	16	5	1	
33959090	a.root-servers.net	1	1	0	
29079969	samba.local.local	5	1	4	
32988913	www.fpu.bg.ac.rs	7	2	0	
18346957	samba.local.local	5	1	4	
26706171	tpc.googlesyndication.com	3	1	0	
15605255	ns.rcub.bg.ac.rs	7	2	1	

27688005	a.root-servers.net	1	1	0
24907918	samba.local.local	5	1	4
17026505	samba.local.local	5	1	4

	w_max	entropy	w_max_ratio	w_count_ratio	digits_ratio	\
24577790	3	2.827820	0.187500	0.062500	0.5625	
33959090	0	-0.000000	0.000000	0.000000	0.0000	
29079969	5	1.921928	1.000000	0.800000	0.0000	
32988913	0	2.128085	0.000000	0.000000	0.0000	
18346957	5	1.921928	1.000000	0.800000	0.0000	
26706171	0	1.584963	0.000000	0.000000	0.0000	
15605255	3	2.807355	0.428571	0.142857	0.0000	
27688005	0	-0.000000	0.000000	0.000000	0.0000	
24907918	5	1.921928	1.000000	0.800000	0.0000	
17026505	5	1.921928	1.000000	0.800000	0.0000	

	uppercase_ratio	time_avg	time_stdev	size_avg	size_stdev	\
24577790	0.0	3.777778	3.270236e+00	18.3	3.301515	
33959090	0.0	31031.111111	3.333333e-01	1.0	0.000000	
29079969	0.0	11.888889	1.105039e+01	5.0	0.000000	
32988913	0.0	989073.555556	1.964106e+06	7.7	0.948683	
18346957	0.0	9.888889	9.157571e+00	5.0	0.000000	
26706171	0.0	11525.000000	1.291751e+04	8.9	11.798776	
15605255	0.0	242450.777778	4.553647e+05	5.4	2.221111	
27688005	0.0	10374.777778	9.619658e+03	1.0	0.000000	
24907918	0.0	8.222222	9.297550e+00	5.0	0.000000	
17026505	0.0	11.111111	1.037358e+01	5.0	0.000000	

	throughput	unique	entropy_avg	entropy_stdev
24577790	5228.571429	0.000000	3.323672	0.365614
33959090	0.035806	1.000000	0.000000	0.000000
29079969	462.962963	1.000000	1.921928	0.000000
32988913	0.008650	0.777778	2.358491	0.365534
18346957	555.555556	1.000000	1.921928	0.000000
26706171	0.858030	0.777778	2.303968	0.781771
15605255	0.024747	0.444444	2.274853	0.526870
27688005	0.107096	1.000000	0.000000	0.000000
24907918	666.666667	1.000000	1.921928	0.000000

Printing the first 10 attacks

```
attack_rows = data[data['attack'] == True].head(10)
print(attack_rows)
```

	user_ip	domain	timestamp	attack	\
3942	186.169.146.147	e5.sk	1624438294225	True	
4297	186.169.146.147	e5.sk	1624438295586	True	
4590	186.169.146.147	e5.sk	1624438296656	True	
6096	186.169.127.58	e5.sk	1624438302237	True	
6187	186.169.146.147	e5.sk	1624438302672	True	
6495	186.169.127.58	e5.sk	1624438303710	True	
6724	186.169.127.58	e5.sk	1624438304691	True	
6968	186.169.127.58	e5.sk	1624438305372	True	
7721	186.169.127.58	e5.sk	1624438308174	True	
7722	186.169.127.58	e5.sk	1624438308181	True	

	request	len	subdomains_count	\
3942	sebubx76xk4erpp3rwehoo3ubmbqeaqbaeaq.a.e.e5.sk	40	3	
4297	4az3kiecotwu3okbtvfm7pdpqcabqeaqbaeaq.a.e.e5.sk	40	3	
4590	x3i2wbqsiucuviqyfaaoxz3lzybqeaqbaeaq.a.e.e5.sk	40	3	
6096	ez2vzwchw3ce5m6wz6cw3nnc2ibqeaqbaeaq.a.e.e5.sk	40	3	
6187	htm7xrligq2enc4lsjhkznd6mbqeaqbaeaq.a.e.e5.sk	40	3	
6495	f4clwtzqaonejfevfnc3vnm334bqeaqbaeaq.a.e.e5.sk	40	3	
6724	hshm7dgsfuvungjbsgjocfazoibqeaqbaeaq.a.e.e5.sk	40	3	
6968	uk7xg4v2usyupazkwfjietmf3ybqeaqbaeaq.a.e.e5.sk	40	3	
7721	ijjuunvalweehk2jgbquu2atwabqeaqbaeaq.a.e.e5.sk	40	3	
7722	mnwmw2m3timeblpdxzjqmnvf3ibqeaqbaeaq.a.e.e5.sk	40	3	

	w_count	w_max	entropy	w_max_ratio	w_count_ratio	digits_ratio	\
3942	3	3	3.975071	0.075	0.075	0.125	
4297	5	3	4.146439	0.075	0.125	0.100	
4590	1	3	3.987326	0.075	0.025	0.075	
6096	1	3	3.893943	0.075	0.025	0.175	
6187	3	3	4.371928	0.075	0.075	0.100	
6495	2	3	3.934830	0.075	0.050	0.125	
6724	3	3	4.137326	0.075	0.075	0.025	
6968	2	3	4.343943	0.075	0.050	0.100	
7721	6	3	3.752656	0.075	0.150	0.050	
7722	3	4	4.062815	0.100	0.075	0.075	

	uppercase_ratio	time_avg	time_stdev	size_avg	size_stdev	\
3942	0.0	2197.222222	2875.261022	48.2	53.370404	
4297	0.0	2348.444444	2779.448601	48.2	53.370404	
4590	0.0	2460.111111	2695.151964	51.8	51.228898	
6096	0.0	1799.222222	1935.781934	44.0	27.712813	
6187	0.0	3105.444444	2782.422466	51.8	51.228898	
6495	0.0	1382.111111	1447.797417	44.0	27.712813	
6724	0.0	1327.555556	1453.227538	44.0	27.712813	
6968	0.0	852.555556	514.388256	44.0	27.712813	
7721	0.0	1163.888889	734.667347	47.6	24.033310	
7722	0.0	1006.777778	819.147389	40.0	0.000000	

	throughput	unique	entropy_avg	entropy_stdev
3942	24.372977	0.0	3.691242	0.910175
4297	22.803615	0.0	3.685581	0.906808
4590	23.394454	0.0	3.884313	0.687639
6096	27.170557	0.0	3.835620	0.663023
6187	18.533095	0.0	3.905225	0.700116
6495	35.369775	0.0	3.824709	0.660105
6724	36.823165	0.0	3.813797	0.653225
6968	57.336461	0.0	3.861596	0.674604
7721	45.437190	0.0	4.036861	0.192920

1. Training on the data provided : only "single" features

1.1 Splitting the data

```
shuffled_data_0 = data.sample(frac=1).reset_index(drop=True)
```

```
shuffled_data_0.head(10)
```

	user_ip	domain	timestamp	attack	request	len	subdomains_count	w_count	w_max	entropy	w_l
0	186.169.145.58	ampproject.org	1624444275435	False	cdn-content.ampproject.org	11	1	7	7	2.663533	
1	130.9.48.79	telephony.goog	1624521121964	False	rs-mts.rcs.telephony.goog	10	2	1	3	2.646439	
2	186.169.123.159	kaspersky-labs.com	1624523977988	False	ksn-crypto-info-geo.kaspersky-labs.com	19	1	6	6	3.642150	
3	186.169.114.79	pki.goog	1624448889645	False	ocsp.pki.goog	4	1	1	3	2.000000	
4	130.9.40.66	local.local	1624515813306	False	samba.local.local	5	1	4	5	1.921928	
5	186.169.4.59	bg.ac.rs	1624523282263	False	proxy.rcub.bg.ac.rs	10	2	7	5	3.121928	
6	130.9.40.66	local.local	1624451279934	False	samba.local.local	5	1	4	5	1.921928	
7	130.9.38.92	root-servers.net	1624503727627	False	a.root-servers.net	1	1	0	0	-0.000000	
8	233.132.83.180	mikrotik.com	1624453303726	False	cloud.mikrotik.com	5	1	6	5	2.321928	
9	130.9.48.195	msftncsi.com	1624493617802	False	dns.msftncsi.com	3	1	0	0	1.584963	

```
shuffled_data_0[shuffled_data_0['attack'] == True].head(10)
```

	user_ip	domain	timestamp	attack
201	199.177.247.90	e5.sk	1624447625053	True
305	186.169.150.60	e5.sk	1624457178293	True
414	186.169.175.72	e5.sk	1624459480374	True
574	186.169.146.147	e5.sk	1624450213304	True
590	186.169.127.58	e5.sk	1624528188340	True
722	186.169.146.147	e5.sk	1624529184853	True
729	199.177.247.53	e5.sk	1624442338674	True
1131	186.169.4.59	e5.sk	1624530694898	True
1260	186.169.146.147	e5.sk	1624450503676	True
1378	186.169.139.60	e5.sk	1624521946272	True

```
# variables
x0 = shuffled_data_0[['len', 'subdomains_count', 'w_count', 'w_count_ratio', 'w_max', 'w_max_ratio', 'digits_ratio', 'entropy']]
```

```
# target
y0 = shuffled_data_0['attack']
```

```
# splitting the data
x_train0, x_test0, y_train0, y_test0 = train_test_split(x0, y0, train_size=0.8)
```

```
print(x_train0.shape)
print(x_test0.shape)
print(y_train0.shape)
print(y_test0.shape)
```

```
(28059320, 8)
(7014830, 8)
(28059320,)
(7014830,)
```

```
print(y_train0.value_counts())
print(y_test0.value_counts())
```

```
attack
False    27919212
True      140108
Name: count, dtype: int64
attack
False     6980159
True       34671
Name: count, dtype: int64
```

```
y_test0.head()
```

```
29283865    False
19184558    False
27011927    False
31828242    False
25171135    False
Name: attack, dtype: bool
```

Même ratio de attack/overall dans le train quand dans le test (0.5%)

```
x_train0.head()
```

```
len  subdomains_count  w_count  w_count_ratio  w_max  w_max_ratio  digits_ratio  entropy
8721718    10              1      11      1.100000      5      0.500000      0.00  3.121928
1533542     5              1       0      0.000000      0      0.000000      0.00  2.321928
14560664    4              1       0      0.000000      0      0.000000      0.50  2.000000
33411955   20              5       4      0.200000      4      0.200000      0.15  3.184184
4161846    21              2      18      0.857143      8      0.380952      0.00  3.689704
```

```
y_train0.head()
```

```
8721718    False
1533542    False
14560664    False
33411955    False
4161846    False
Name: attack, dtype: bool
```

✓ 1.2 The first model : a Logistic Regression

```
model01 = LogisticRegression(max_iter=10000)
model01.fit(x_train0, y_train0)
```

```
LogisticRegression
LogisticRegression(max_iter=10000)
```

```
train_predictions = model01.predict(x_train0)
train_score = accuracy_score(y_train0, train_predictions)
print(f"Training accuracy : {train_score}")
recall = recall_score(y_train0, train_predictions)
print(f"Recall score: {recall}")
print(confusion_matrix(y_train0, train_predictions))
```

```
Training accuracy : 0.9977963471673583
Recall score: 0.7226639449567477
[[27896236  22976]
 [  38857  101251]]
```

```

predictions = model01.predict(x_test0)
test_score = accuracy_score(y_test0, predictions)
print(f"Testing accuracy : {test_score}")
recall = recall_score(y_test0, predictions)
print(f"Recall score: {recall_score}")
print(confusion_matrix(y_test0,predictions))

```

Testing accuracy : 0.9978073595511224
Recall score: <function recall_score at 0x7cd27966d990>
[[6974333 5826]
[9555 25116]]

Vraiment pas terrible. C'est d'ailleurs les résultats du papier, donc pas de grosse surprise. L'accuracy est vraiment pas un bon représentant de la performance étant donnée la petite proportion de "true". Là, on détecte 72% des attaques

1.3 The second model: a MLP

```

# Define the model architecture
model02 = tf.keras.Sequential([
    tf.keras.layers.Dense(64, activation='relu', input_shape=(x_train0.shape[1],)),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

# Compile the model
model02.compile(optimizer='adam',
                loss='binary_crossentropy',
                metrics=['accuracy', tf.keras.metrics.Precision()])

# Print the model summary
model02.summary()

```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
dense_9 (Dense)	(None, 64)	576
dropout_6 (Dropout)	(None, 64)	0
dense_10 (Dense)	(None, 64)	4160
dropout_7 (Dropout)	(None, 64)	0
dense_11 (Dense)	(None, 1)	65
Total params: 4801 (18.75 KB)		
Trainable params: 4801 (18.75 KB)		
Non-trainable params: 0 (0.00 Byte)		

```

# Train the model
history = model02.fit(x_train0, y_train0, epochs=4, batch_size=2048, validation_data=(x_test0, y_test0))

```

Epoch 1/4
13701/13701 [=====] - 79s 6ms/step - loss: 0.0080 - accuracy: 0.9981 - precision_3: 0.7928 - val_loss: 0.0080 - val_accuracy: 0.9981 - val_precision_3: 0.7928
Epoch 2/4
13701/13701 [=====] - 75s 5ms/step - loss: 0.0023 - accuracy: 0.9996 - precision_3: 0.9673 - val_loss: 0.0023 - val_accuracy: 0.9996 - val_precision_3: 0.9673
Epoch 3/4
13701/13701 [=====] - 74s 5ms/step - loss: 0.0021 - accuracy: 0.9996 - precision_3: 0.9707 - val_loss: 0.0021 - val_accuracy: 0.9996 - val_precision_3: 0.9707
Epoch 4/4
13701/13701 [=====] - 75s 6ms/step - loss: 0.0020 - accuracy: 0.9996 - precision_3: 0.9744 - val_loss: 0.0020 - val_accuracy: 0.9996 - val_precision_3: 0.9744

```

# Make predictions on the test set
y_pred = model02.predict(x_test0)
y_pred_classes = (y_pred > 0.5).astype("int32") # Assuming a binary classification with a threshold of 0.5

```

219214/219214 [=====] - 266s 1ms/step

```

# Compute the confusion matrix
cm = confusion_matrix(y_test0, y_pred_classes)
print(cm)

```

[[6979991 168]
[1538 33133]]

Super résultats !!!! C'est quasi parfait sur les false, et c'est pas si mal sur les true : 96% des attaques ont été détectées !

✓ 2. Training with data from scratch

```
# Going from scratch  
data_ini = data[['request', 'attack']]
```

✓ 2.1 Preprocessing

```

# Calculate the length of the request (excluding the TLD)
def calculate_len(request):
    parts = request.split('.')
    tld_length = len(parts[-1]) + len(parts[-2])
    return len(request) - tld_length - 2 # exclude the 2 dots

# Calculate number of subdomains
def calculate_subdomains(request):
    return request.count('.') - 1

# Count the number of English words in the request
def calculate_w_count(request):
    delimiters = ['.', '-', '_']
    words = request
    for delimiter in delimiters:
        words = words.replace(delimiter, ' ')
    words = words.split()
    #print(words)

    count = 0
    for word in words:
        if word.lower() in english_words:
            count += 1
            #print(word)
    return count

# Calculate the length of the longest English word in the request
def calculate_longest_word_length(request):
    delimiters = ['.', '-', '_']
    words = request
    for delimiter in delimiters:
        words = words.replace(delimiter, ' ')
    words = words.split()

    longest_length = 0
    for word in words:
        if word.lower() in english_words:
            if len(word) > longest_length:
                longest_length = len(word)

    return longest_length

# Calculate the percentage of digits in the request
def count_digits(text):
    return sum(char.isdigit() for char in text)

# Calculate the entropy of the request
def calculate_entropy(request):
    probability = [float(request.count(c)) / len(request) for c in dict.fromkeys(list(request))]
    entropy = - sum([p * math.log(p) / math.log(2.0) for p in probability])
    return entropy

# Function to calculate entropy of a string
def calculate_entropy(text):
    if len(text) == 0:
        return 0.0

    # Calculate frequency of each character
    freq = {}
    for char in text:
        if char in freq:
            freq[char] += 1
        else:
            freq[char] = 1

    # Calculate entropy
    entropy = 0.0
    text_length = len(text)
    for count in freq.values():
        probability = count / text_length
        entropy -= probability * math.log2(probability)

    return entropy

data_test=data_ini.head(30)

```

```
data_test['len'] = data_test['request'].apply(calculate_len)
data_test['subdomains_count'] = data_test['request'].apply(calculate_subdomains)
data_test['w_count'] = data_test['request'].apply(calculate_w_count)
data_test['w_count_ratio'] = data_test['w_count'] / data_test['len']
data_test['w_max'] = data_test['request'].apply(calculate_longest_word_length)
data_test['w_max_ratio'] = data_test['w_max'] / data_test['len']
data_test['digit_count'] = data_test['request'].apply(count_digits)
data_test['digit_ratio'] = data_test['digit_count'] / data_test['len']
data_test['entropy'] = data_test['request'].apply(calculate_entropy)

data_test.head(30)
```

```

<ipython-input-113-adc5d2d8e2c1>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-wrapper-around-data
data_test['len'] = data_test['request'].apply(calculate_len)
<ipython-input-113-adc5d2d8e2c1>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-wrapper-around-data
data_test['subdomains_count'] = data_test['request'].apply(calculate_subdomains)
<ipython-input-113-adc5d2d8e2c1>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-wrapper-around-data
data_test['w_count'] = data_test['request'].apply(calculate_w_count)
<ipython-input-113-adc5d2d8e2c1>:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-wrapper-around-data
data_test['w_count_ratio'] = data_test['w_count'] / data_test['len']
<ipython-input-113-adc5d2d8e2c1>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-wrapper-around-data
data_test['w_max'] = data_test['request'].apply(calculate_longest_word_length)
<ipython-input-113-adc5d2d8e2c1>:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-wrapper-around-data
data_test['w_max_ratio'] = data_test['w_max'] / data_test['len']
<ipython-input-113-adc5d2d8e2c1>:7: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-wrapper-around-data
data_test['digit_count'] = data_test['request'].apply(count_digits)
<ipython-input-113-adc5d2d8e2c1>:8: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-wrapper-around-data
data_test['digit_ratio'] = data_test['digit_count'] / data_test['len']
<ipython-input-113-adc5d2d8e2c1>:9: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-wrapper-around-data
data_test['entropy'] = data_test['request'].apply(calculate_entropy)

```

	request	attack	len	subdomains_count	w_count	w_count_ratio	w_max	w_max_ratio	digit_count	d
0	f.surbl.org	False	1	1	0	0.000000	0	0.000000	0	
1	118.141.11.106.sbl.spamhaus.org	False	18	5	0	0.000000	0	0.000000	11	
2	118.141.11.106.zen.spamhaus.org	False	18	5	0	0.000000	0	0.000000	11	
3	128.141.11.106.sbl.spamhaus.org	False	18	5	0	0.000000	0	0.000000	11	
4	128.141.11.106.zen.spamhaus.org	False	18	5	0	0.000000	0	0.000000	11	
5	68.211.11.106.sbl.spamhaus.org	False	17	5	0	0.000000	0	0.000000	10	
6	68.211.11.106.zen.spamhaus.org	False	17	5	0	0.000000	0	0.000000	10	
7	28.41.205.140.sbl.spamhaus.org	False	17	5	0	0.000000	0	0.000000	10	
8	28.41.205.140.zen.spamhaus.org	False	17	5	0	0.000000	0	0.000000	10	
9	17.41.205.140.sbl.spamhaus.org	False	17	5	0	0.000000	0	0.000000	10	
10	17.41.205.140.zen.spamhaus.org	False	17	5	0	0.000000	0	0.000000	10	
11	127.141.11.106.sbl.spamhaus.org	False	18	5	0	0.000000	0	0.000000	11	
12	127.141.11.106.zen.spamhaus.org	False	18	5	0	0.000000	0	0.000000	11	
13	67.211.11.106.sbl.spamhaus.org	False	17	5	0	0.000000	0	0.000000	10	
14	67.211.11.106.zen.spamhaus.org	False	17	5	0	0.000000	0	0.000000	10	
15	57.211.11.106.sbl.spamhaus.org	False	17	5	0	0.000000	0	0.000000	10	
16	57.211.11.106.zen.spamhaus.org	False	17	5	0	0.000000	0	0.000000	10	
17	27.41.205.140.sbl.spamhaus.org	False	17	5	0	0.000000	0	0.000000	10	
18	27.41.205.140.zen.spamhaus.org	False	17	5	0	0.000000	0	0.000000	10	

19	17.81.205.140.sbl.spamhaus.org	False	17	5	0	0.000000	0	0.000000	10
20	17.81.205.140.zen.spamhaus.org	False	17	5	0	0.000000	0	0.000000	10
21	117.141.11.106.sbl.spamhaus.org	False	18	5	0	0.000000	0	0.000000	11
22	117.141.11.106.zen.spamhaus.org	False	18	5	0	0.000000	0	0.000000	11
23	27.81.205.140.sbl.spamhaus.org	False	17	5	0	0.000000	0	0.000000	10
24	27.81.205.140.zen.spamhaus.org	False	17	5	0	0.000000	0	0.000000	10
25	20.177.89.47.sbl.spamhaus.org	False	16	5	0	0.000000	0	0.000000	9
26	20.177.89.47.zen.spamhaus.org	False	16	5	0	0.000000	0	0.000000	9
27	ksn-crypto-url-geo.kas-labs.com	False	18	1	1	0.055556	4	0.222222	0
28	ksn-crypto-url-geo.kaspersky-labs.com	False	18	1	1	0.055556	4	0.222222	0

```
# Pour comparer avec les features initiales
data.head(30)
```


	user_ip	domain	timestamp	attack	request	len	subdomains_count	w_count	w_max
0	186.169.253.58	surbl.org	1624438272607	False	f.surbl.org	1	1	0	0
1	186.169.253.58	spamhaus.org	1624438273058	False	118.141.11.106.sbl.spamhaus.org	18	5	0	0
2	186.169.253.58	spamhaus.org	1624438273058	False	118.141.11.106.zen.spamhaus.org	18	5	1	3
3	186.169.253.58	spamhaus.org	1624438273059	False	128.141.11.106.sbl.spamhaus.org	18	5	0	0
4	186.169.253.58	spamhaus.org	1624438273059	False	128.141.11.106.zen.spamhaus.org	18	5	1	3
5	186.169.253.58	spamhaus.org	1624438273060	False	68.211.11.106.sbl.spamhaus.org	17	5	0	0
6	186.169.253.58	spamhaus.org	1624438273060	False	68.211.11.106.zen.spamhaus.org	17	5	1	3
7	186.169.253.58	spamhaus.org	1624438273060	False	28.41.205.140.sbl.spamhaus.org	17	5	0	0
8	186.169.253.58	spamhaus.org	1624438273061	False	28.41.205.140.zen.spamhaus.org	17	5	1	3
9	186.169.253.58	spamhaus.org	1624438273062	False	17.41.205.140.sbl.spamhaus.org	17	5	0	0
10	186.169.253.58	spamhaus.org	1624438273062	False	17.41.205.140.zen.spamhaus.org	17	5	1	3
11	186.169.253.58	spamhaus.org	1624438273063	False	127.141.11.106.sbl.spamhaus.org	18	5	0	0
12	186.169.253.58	spamhaus.org	1624438273063	False	127.141.11.106.zen.spamhaus.org	18	5	1	3
13	186.169.253.58	spamhaus.org	1624438273063	False	67.211.11.106.sbl.spamhaus.org	17	5	0	0
14	186.169.253.58	spamhaus.org	1624438273064	False	67.211.11.106.zen.spamhaus.org	17	5	1	3
15	186.169.253.58	spamhaus.org	1624438273064	False	57.211.11.106.sbl.spamhaus.org	17	5	0	0
16	186.169.253.58	spamhaus.org	1624438273064	False	57.211.11.106.zen.spamhaus.org	17	5	1	3
17	186.169.253.58	spamhaus.org	1624438273065	False	27.41.205.140.sbl.spamhaus.org	17	5	0	0
18	186.169.253.58	spamhaus.org	1624438273065	False	27.41.205.140.zen.spamhaus.org	17	5	1	3
19	186.169.253.58	spamhaus.org	1624438273066	False	17.81.205.140.sbl.spamhaus.org	17	5	0	0
20	186.169.253.58	spamhaus.org	1624438273066	False	17.81.205.140.zen.spamhaus.org	17	5	1	3
21	186.169.253.58	spamhaus.org	1624438273066	False	117.141.11.106.sbl.spamhaus.org	18	5	0	0
22	186.169.253.58	spamhaus.org	1624438273067	False	117.141.11.106.zen.spamhaus.org	18	5	1	3
23	186.169.253.58	spamhaus.org	1624438273067	False	27.81.205.140.sbl.spamhaus.org	17	5	0	0
24	186.169.253.58	spamhaus.org	1624438273067	False	27.81.205.140.zen.spamhaus.org	17	5	1	3
25	186.169.253.58	spamhaus.org	1624438273100	False	20.177.89.47.sbl.spamhaus.org	16	5	0	0
26	186.169.253.58	spamhaus.org	1624438273101	False	20.177.89.47.zen.spamhaus.org	16	5	1	3
27	186.169.123.159	kas-labs.com	1624438273201	False	ksn-crypto-url-geo.kas-labs.com	18	1	4	6
28	186.169.123.159	kaspersky-labs.com	1624438273204	False	ksn-crypto-url-geo.kaspersky-labs.com	18	1	4	6

```
data_1 = data_ini
```

```

data_1['len'] = data_1['request'].apply(calculate_len)
data_1['subdomains_count'] = data_1['request'].apply(calculate_subdomains)
data_1['w_count'] = data_1['request'].apply(calculate_w_count)
data_1['w_count_ratio'] = data_1['w_count'] / data_1['len']
data_1['w_max'] = data_1['request'].apply(calculate_longest_word_length)
data_1['w_max_ratio'] = data_1['w_max'] / data_1['len']
data_1['digit_count'] = data_1['request'].apply(count_digits)
data_1['digit_ratio'] = data_1['digit_count'] / data_1['len']
data_1['entropy'] = data_1['request'].apply(calculate_entropy)

```

 <ipython-input-116-ce7c54dcc958>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a

```

data_1['len'] = data_1['request'].apply(calculate_len)
<ipython-input-116-ce7c54dcc958>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a

```

data_1['subdomains_count'] = data_1['request'].apply(calculate_subdomains)
<ipython-input-116-ce7c54dcc958>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a

```

data_1['w_count'] = data_1['request'].apply(calculate_w_count)
<ipython-input-116-ce7c54dcc958>:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a

```

data_1['w_count_ratio'] = data_1['w_count'] / data_1['len']
<ipython-input-116-ce7c54dcc958>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a

```

data_1['w_max'] = data_1['request'].apply(calculate_longest_word_length)
<ipython-input-116-ce7c54dcc958>:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a

```

data_1['w_max_ratio'] = data_1['w_max'] / data_1['len']
<ipython-input-116-ce7c54dcc958>:7: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a

```

data_1['digit_count'] = data_1['request'].apply(count_digits)
<ipython-input-116-ce7c54dcc958>:8: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a

```

data_1['digit_ratio'] = data_1['digit_count'] / data_1['len']
<ipython-input-116-ce7c54dcc958>:9: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a

```

data_1['entropy'] = data_1['request'].apply(calculate_entropy)

```

data_1.info()


 <class 'pandas.core.frame.DataFrame'>
RangeIndex: 35074150 entries, 0 to 35074149
Data columns (total 11 columns):
Column Dtype

0 request object
1 attack bool
2 len int64
3 subdomains_count int64
4 w_count int64
5 w_count_ratio float64
6 w_max int64
7 w_max_ratio float64
8 digit_count int64
9 digit_ratio float64
10 entropy float64
dtypes: bool(1), float64(4), int64(5), object(1)
memory usage: 2.6+ GB

2.2 Splitting the data

```
shuffled_data_1 = data_1.sample(frac=1).reset_index(drop=True)
```

```
shuffled_data_1.head(10)
```



	request	attack	len	subdomains_count	w_count	w_count_ratio	w_max	w_max_ratio	digit_count	digit_ra
0	i-43.b-46010.ut.bench.utorrent.com	False	21	4	1	0.047619	5	0.238095	7	0.333333
1	samba.local.local	False	5	1	3	0.600000	5	1.000000	0	0.000000
2	mail.ffh.bg.ac.rs	False	11	3	1	0.090909	4	0.363636	0	0.000000
3	web.facebook.com	False	3	1	1	0.333333	3	1.000000	0	0.000000
4	www.google.com	False	3	1	0	0.000000	0	0.000000	0	0.000000
5	samba.local.local	False	5	1	3	0.600000	5	1.000000	0	0.000000
6	teams-events-data.trafficmanager.net	False	17	1	4	0.235294	6	0.352941	0	0.000000
7	samba.local.local	False	5	1	3	0.600000	5	1.000000	0	0.000000
8	ksn-cinfo-geo.kas-labs.com	False	13	1	1	0.076923	4	0.307692	0	0.000000

```
shuffled_data_1[shuffled_data_1['attack'] == True].head(10)
```



	request	attack	1
95	efysfpxsydxufcp7tihtyse5gebqeaqbaeaq.a.e.e5.sk	True	
255	5lw5dg6ythiuja7egdpgsr2tjibqeaqbaeaq.a.e.e5.sk	True	
480	AlXsV_4iZR.dnsresearch.ml	True	
700	ScrQtP1j6MfXEGsF82KxzlVBO4OOuaX9EFZ-9W5s2UbVHwCMY5D-XTb0WpQJaW.oEMG5RB.dnsresearch.ml	True	
1363	ag6cuaacaakityrgaeaaaaaaaaabveinbki3cpceyqqm3k7aeaaaaa3aaaaa3s.yomnacaawaajqb57r7wr26nlopvkqryhstj2mxsrgefzs2ai.a.j.e5.sk	True	
2311	n764ciz43evu7prz7dfkga2744bqeaqbaeaq.a.e.e5.sk	True	
2341	hyesuaacaakjngaaaaaeaaaaaaaaabr5u4g2ml4pefekf5acaeeaaaaa3aaabus.aktvacaawaajqbib4wdfwx2btrccv4c5hpozlrmsamqm3aai.a.j.e5.sk	True	
2625	iluTZ0bfyQV-gGhDcxBN7XlvApM8De.m4FEt9cR-GII_-EEq0Wuzsod5Jh751.izV65jgwi.dnsresearch.ml	True	
2794	ixdo7aoqgeburlycyqrphflombqeaqbaeaq.a.e.e5.sk	True	
2815	o76n2dfeldzuzni44dmvbhe5xmbqeaqbaeaq.a.e.e5.sk	True	

```
# variables
x1 = shuffled_data_1[['len','subdomains_count','w_count','w_count_ratio','w_max','w_max_ratio', 'digit_ratio','entropy']]

# target
y1 = shuffled_data_1['attack']

# splitting the data
x_train1, x_test1, y_train1, y_test1 = train_test_split(x1, y1, train_size=0.8)

print(x_train1.shape)
print(x_test1.shape)
print(y_train1.shape)
print(y_test1.shape)

(28059320, 8)
(7014830, 8)
(28059320,)
(7014830,)

print(y_train1.value_counts())
print(y_test1.value_counts())

attack
False    27919538
True      139782
Name: count, dtype: int64
attack
False     6979833
True       34997
Name: count, dtype: int64
```

```
# Pour obtenir des queries du test set
false_indexes = y_test1[y_test1 == False].index

ok_in_test = shuffled_data_1.loc[false_indexes]

ok_in_test.head()
```

	request	attack	len	subdomains_count	w_count	w_count_ratio	w_max	w_max_ratio	digit_count	digit_
21393812	star.c10r.facebook.com	False	9	2	1	0.111111	4	0.444444	2	0.
27325187	ssl.gstatic.com	False	3	1	0	0.000000	0	0.000000	0	0.
4858532	sirius.mwbsys.com	False	6	1	0	0.000000	0	0.000000	0	0.
7501300	a2047.r.akamai.net	False	7	2	1	0.142857	3	0.428571	4	0.
24000768	a.root-servers.net	False	1	1	4	4.000000	7	7.000000	0	0.

```
# Pour obtenir des queries du test set
true_indexes = y_test1[y_test1].index

attacks_in_test = shuffled_data_1.loc[true_indexes]

attacks_in_test.head()
```

	request	attack	len	subdomains_count	w_count	w_count_ratio	w_max	w_max_rati
13556679	90dNIhWO8jgmZbsGLT_dnsresearch.ml	True	20	1	0	0.000	0	0.00
20085422	zfw5cdrxwxcehuvnbskuf6jcqbqaeqbaeaq.a.e.e5.sk	True	40	3	1	0.025	1	0.02
21399344	ffaolftgc2euvueuhcrrr3hfoabqaeqbaeaq.a.e.e5.sk	True	40	3	1	0.025	1	0.02
22693012	2hb6r2llewbulpmp2ioizgjhqubqaeqbaeaq.a.e.e5.sk	True	40	3	1	0.025	1	0.02
22654698	7fozyh3ozjmerdknbc5yjb74ybqaeqbaeaq.a.e.e5.sk	True	40	3	1	0.025	1	0.02

2.3 The first model : a Logistic regression

```
modell = LogisticRegression(max_iter=10000)
modell.fit(x_train1, y_train1)
```

```
LogisticRegression
LogisticRegression(max_iter=10000)
```

```
train_predictions = modell.predict(x_train1)
train_score = accuracy_score(y_train1, train_predictions)
recall = recall_score(y_train1, train_predictions)
print(f"Recall score: {recall}")
print(f"Training accuracy : {train_score}")
print(confusion_matrix(y_train1, train_predictions))
```

```
Recall score: 0.6999828304073485
Training accuracy : 0.9979170557233746
[[27903029  16509]
 [ 41937  97845]]
```

```
predictions = modell.predict(x_test1)
test_score = accuracy_score(y_test1, predictions)
print(f"Testing accuracy : {test_score}")
recall = recall_score(y_test1, predictions)
print(f"Recall score: {recall}")
print(confusion_matrix(y_test1, predictions))
```

```
Testing accuracy : 0.9979127077919209
Recall score: 0.6996885447324056
[[6975701  4132]
 [ 10510  24487]]
```

Soit 70% des attaques détectées C'est pas si mal.

2.4 The second model : a MLP


```
# Define the model architecture
model = tf.keras.Sequential([
    tf.keras.layers.Dense(64, activation='relu', input_shape=(x_train1.shape[1],)),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

```
# Compile the model
model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['accuracy', tf.keras.metrics.Precision()])
```

```
# Print the model summary
model.summary()
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
dense_12 (Dense)	(None, 64)	576
dropout_8 (Dropout)	(None, 64)	0
dense_13 (Dense)	(None, 64)	4160
dropout_9 (Dropout)	(None, 64)	0
dense_14 (Dense)	(None, 1)	65
Total params: 4801 (18.75 KB)		
Trainable params: 4801 (18.75 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
# Train the model
history = model.fit(x_train1, y_train1, epochs=4, batch_size=2048, validation_data=(x_test1, y_test1))
```

Epoch 1/4
 13701/13701 [=====] - 79s 6ms/step - loss: 0.0074 - accuracy: 0.9985 - precision_4: 0.8449 - val_loss: 0.0074 - val_accuracy: 0.9985 - val_precision_4: 0.8449
 Epoch 2/4
 13701/13701 [=====] - 76s 6ms/step - loss: 0.0020 - accuracy: 0.9996 - precision_4: 0.9802 - val_loss: 0.0020 - val_accuracy: 0.9996 - val_precision_4: 0.9802
 Epoch 3/4
 13701/13701 [=====] - 76s 6ms/step - loss: 0.0019 - accuracy: 0.9996 - precision_4: 0.9812 - val_loss: 0.0019 - val_accuracy: 0.9996 - val_precision_4: 0.9812
 Epoch 4/4
 13701/13701 [=====] - 75s 5ms/step - loss: 0.0018 - accuracy: 0.9997 - precision_4: 0.9818 - val_loss: 0.0018 - val_accuracy: 0.9997 - val_precision_4: 0.9818

```
# Make predictions on the test set
y_pred = model.predict(x_test1)
y_pred_classes = (y_pred > 0.5).astype("int32") # Assuming a binary classification with a threshold of 0.5
```

219214/219214 [=====] - 262s 1ms/step

```
# Compute the confusion matrix
cm = confusion_matrix(y_test1, y_pred_classes)
print(cm)
```

```
[[6979605    228]
 [      829 34168]]
```

Précision de 92%

2.4 Evulating manual inputs

https://colab.research.google.com/drive/1ctl4H50yofAcxztxnl_UO8xqj3b8Meoa#scrollTo=ZspTBRZlXC4c&printMode=true 18/20

3. Looking for a better model

Le passage de 1. à 2. a montré qu'on s'en sort bien en computant nous mêmes les features (même si y a des différences entre nos features et les features du dataset, ces différences ne sont donc pas significatives).

Maintenant, on a un problème : le modèle, en gros, dit que tout input est FALSE sauf les inputs qui ressemblent vraiment aux attaques du dataset (92% de précision quand même) D'où les résultats observés pour les manual inputs, et du coup notre modèle ne nous sert pas dans l'application

Ca, c'est dû à la quantité relativement très faibles de TRUE attacks dans le dataset. Pour parer à ça, deux approches : utiliser une nouvelle loss fonction faite pour prévoir le coup, ou bien ne retenir que très peu de rows du dataset pour un ratio 1 true pour 1 false

3.1 Focal loss function

```
# Define the model architecture
model3 = tf.keras.Sequential([
    tf.keras.layers.Dense(64, activation='relu', input_shape=(x_train1.shape[1],)),
    tf.keras.layers.Dropout(0.7),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dropout(0.7),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

# Compile the model
model3.compile(optimizer='adam',
               loss=tf.keras.losses.BinaryFocalCrossentropy(apply_class_balancing=True, alpha=0.99, gamma=4.0, from_logits=True),
               metrics=['accuracy', tf.keras.metrics.Precision()])

# Print the model summary
model3.summary()
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
dense_15 (Dense)	(None, 64)	576
dropout_10 (Dropout)	(None, 64)	0
dense_16 (Dense)	(None, 64)	4160
dropout_11 (Dropout)	(None, 64)	0
dense_17 (Dense)	(None, 1)	65
Total params: 4801 (18.75 KB)		
Trainable params: 4801 (18.75 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
# Train the model
history = model3.fit(x_train1, y_train1, epochs=4, batch_size=4096, validation_data=(x_test1, y_test1))
```

```
Epoch 1/4
6851/6851 [=====] - 49s 7ms/step - loss: 7.3814e-04 - accuracy: 0.9744 - precision_5: 0.1550 -
Epoch 2/4
6851/6851 [=====] - 46s 7ms/step - loss: 7.2798e-05 - accuracy: 0.9887 - precision_5: 0.3012 -
Epoch 3/4
6851/6851 [=====] - 46s 7ms/step - loss: 6.7300e-05 - accuracy: 0.9890 - precision_5: 0.3075 -
Epoch 4/4
6851/6851 [=====] - 45s 7ms/step - loss: 6.8769e-05 - accuracy: 0.9879 - precision_5: 0.2874 -
```

```
# Make predictions on the test set
y_pred = model3.predict(x_test1)
y_pred_classes = (y_pred > 0.5).astype("int32") # Assuming a binary classification with a threshold of 0.5
```

```
219214/219214 [=====] - 262s 1ms/step
```

```
# Compute the confusion matrix
cm = confusion_matrix(y_test1, y_pred_classes)
print(cm)
```

```
[[6976190 3643]
 [ 2602 32395]]
```

```
for i in range(len(queries)):
    query = queries[i]
    input = preprocessing(query)
    input = input.reshape(1, -1) # Reshape to add batch dimension
```

```
pred = model3.predict(input)
print("Input query : ", query)
print("It's an attak : ", labels[i])
print("Prediction:", pred, "=> ", pred > 0.3)
```

```
1/1 [=====] - 0s 62ms/step
Input query : mawx0cmf0zwqgb3zlcibetlmuckhlcmugaxmgysbwyxnzd29yzdoguebzc3cwcw.example.com
It's an attak : True
Prediction: [[0.36274105]] => [[ True]]
1/1 [=====] - 0s 25ms/step
Input query : qxmjmhckfuzbozxjliglzigegec2vjcmv0igtletogu0vdukvus0vzmtizndu2.example.com
It's an attak : True
Prediction: [[0.36149856]] => [[ True]]
1/1 [=====] - 0s 23ms/step
Input query : vghpcybpbyzb2llihnlbnnpdgl2zsbkxrhroyxqgbmvzlzhmgdggymugzxh.example.com
It's an attak : True
Prediction: [[0.36291838]] => [[ True]]
1/1 [=====] - 0s 24ms/step
Input query : samba.local.local
It's an attak : False
Prediction: [[0.01384288]] => [[False]]
1/1 [=====] - 0s 23ms/step
Input query : a.c-0.19-a3000000.d0c0081.1838.1220.2fc9.410.0.kutu452468r6pmrknpzkt6lt.avqs.mcafee.com
It's an attak : True
Prediction: [[0.8810738]] => [[ True]]
1/1 [=====] - 0s 25ms/step
Input query : m27suaacaakc2obqahqeya6saqaaa4chqaeviljygcpxgaaaaa3aaaabu6.vj7nacaawaajqaro5ttfstvahphtk3asngfy5k7
It's an attak : True
Prediction: [[0.98675406]] => [[ True]]
1/1 [=====] - 0s 26ms/step
Input query : ihccuaacaakpcyjha4gksaqaanawdsq7f4lbe4r7bjtqaeaaaaa3aaaabv5.4lmaacaawaajqb62bnxydj5qus24g7y4bh6icer
It's an attak : True
Prediction: [[0.9858124]] => [[ True]]
1/1 [=====] - 0s 27ms/step
Input query : colab.research.google.com
It's an attak : False
Prediction: [[0.03670552]] => [[False]]
1/1 [=====] - 0s 24ms/step
Input query : colab.researsh.google.com
It's an attak : True
Prediction: [[0.33043373]] => [[ True]]
```

Je trouve ça pas mal comme résultat, car si on change le threshold à 0.30, on obtient une bonne prédiction Remarque :

"colab.research.google.com" est déclaré false, mais "colab.researsh.google.com" est déclaré true => impact des mots. Une faute de frappe ou d'orthographe a une grosse conséquence

✓ 3.2 Looking for the threshold appropriate to our application

J'ai réalisé une exfiltration avec un grand fichier (diary.txt) afin d'obtenir de nombreuses query générées par mes scripts. Le but est d'obtenir un set pertinent pour notre application, afin de voir les prédictions du modèles dessus, puis de déterminer un threshold adapté.

```
# Load the sample of "real" attack queries
with open(os.path.join(folder_path, 'queries.txt'), 'r') as f:
    lines = f.readlines()
    line = [line.strip() for line in lines]

real_att_queries = []
for i in range(len(line)):
    if i%2 == 0 :
        real_att_queries.append(line[i])

print(len(real_att_queries))
print(real_att_queries)
```

```
583
['tg9yzw0gaxbzd0gzg9sb3igc2l0igftzxsiggnvbnly3rldhvyigfkaxbpc2n.example.com.', 'pbmkgwxpdc4guhjvaw4gy29tbw9kbybzzwqgt
```