

Point d'avancement PRIM Exploitation de Cilium et Hubble pour détecter et se protéger d'attaques DNS exfiltration

30 Mai 2024

Gilles HOPIN

Professeurs : Jean-Louis ROUGIER, Patrice NIVAGGIOLI

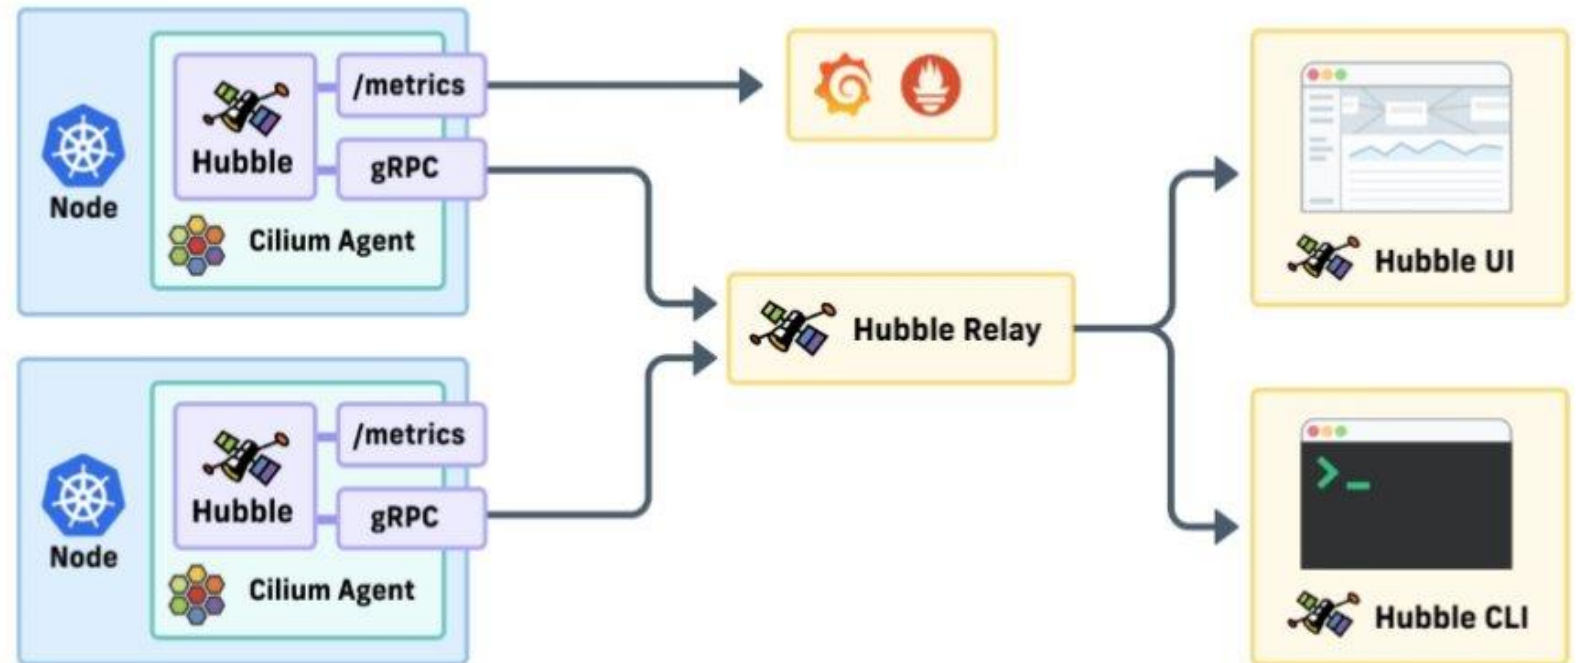
Sommaire

- Recap sur ce qu'offre hubble : metrics et logs
- Exportation des network flow logs
 - Configuration d'Hubble exporter
 - Configuration de Loki
 - Problème rencontré avec Promtail
 - Récupération des logs pour classification
- Classification => un problème de ML
 - Article Splunk : leur approche
 - Un dataset trouvé
 - Perspectives ML
- To do list

Ce qu'offre Hubble


Hubble permet de faire remonter deux types de données :

- Des metrics
- Des network flow logs




Les metrics



Une
personnalis
ation assez
limitée...

 Prometheus Alerts Graph Status ▾ Help

☐ Use local time ☐ Enable query history ☒ Enable autocomplete ☒ Enable highlighting ☒ Enable linter



hubble_dns_queries_total



Execute

Load time: 106ms Resolution: 14s Result series: 3

Table

Graph

<

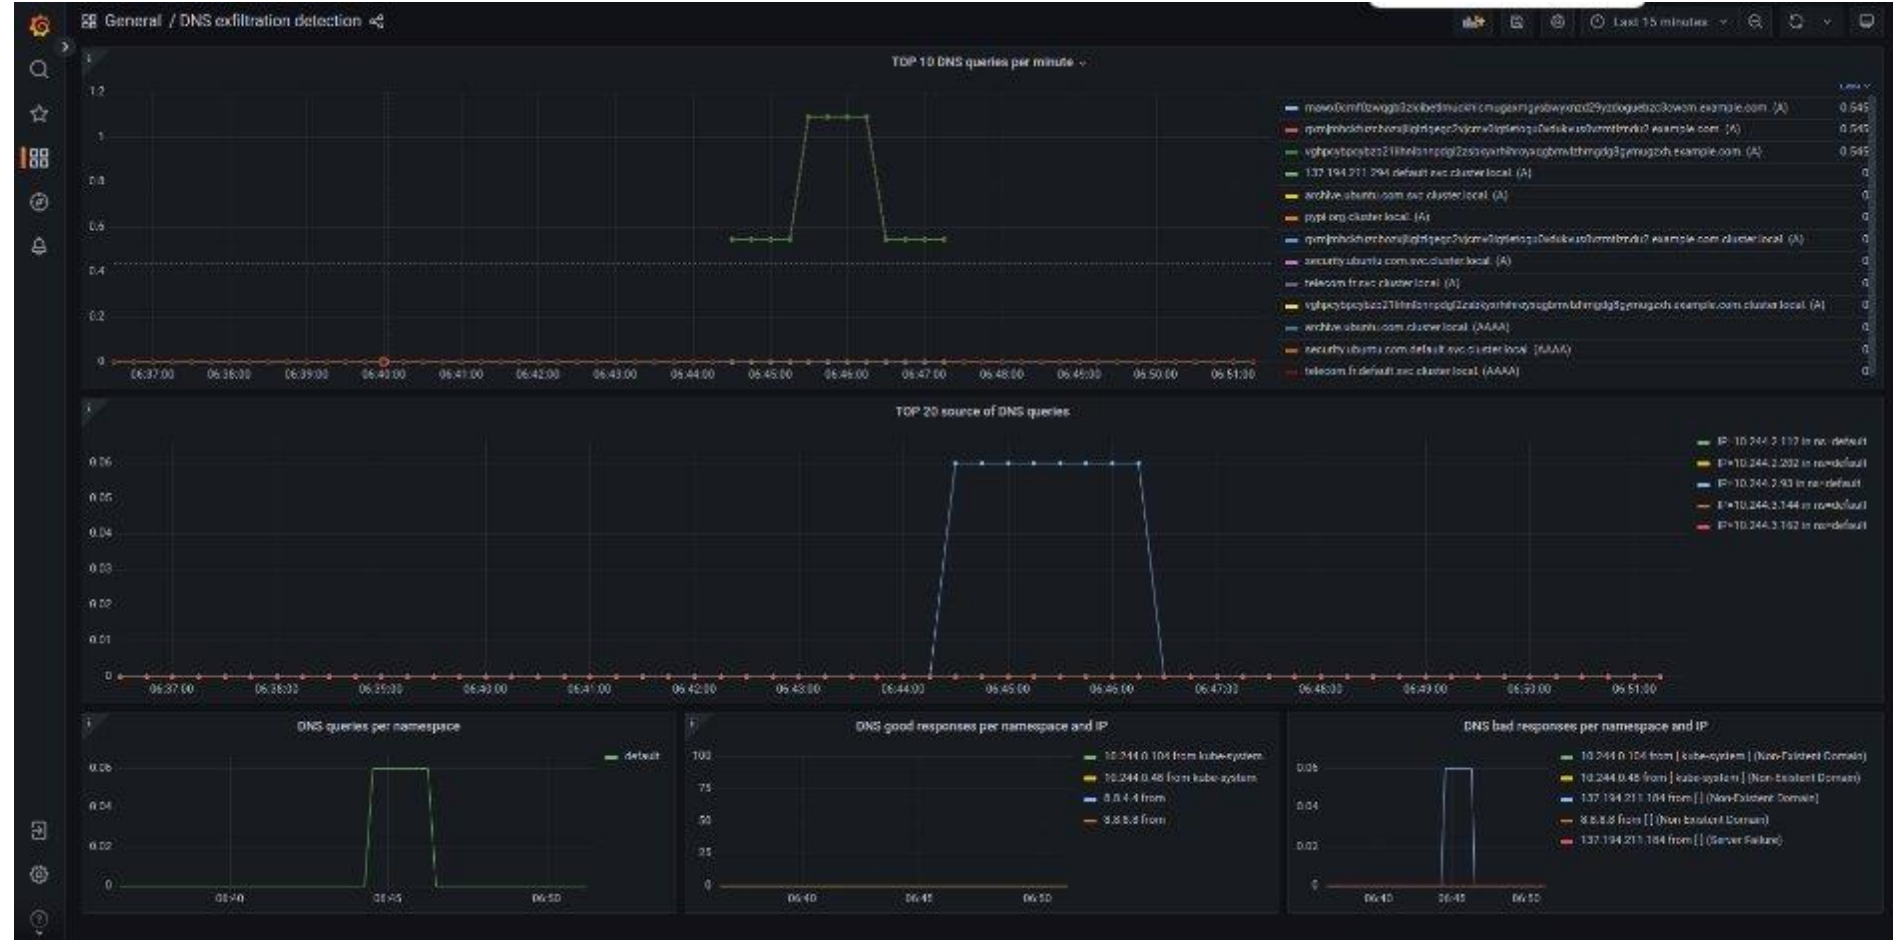
Evaluation time

>

hubble_dns_queries_total{app_kubernetes_io_managed_by="Helm", app_kubernetes_io_name="hubble", app_kubernetes_io_part_of="cilium", instance="172.18.0.3:9965", ips_returned="0", job="kubernetes-service-endpoints", k8s_app="hubble", namespace="kube-system", node="kind-worker3", qtypes="A", query="mawx0cmf0zwqgb3zlcibetlmuckhlcmugaxmgysbwynzd29yzdoguebzc3cwcm.example.com.", service="hubble-metrics"}	43
hubble_dns_queries_total{app_kubernetes_io_managed_by="Helm", app_kubernetes_io_name="hubble", app_kubernetes_io_part_of="cilium", instance="172.18.0.3:9965", ips_returned="0", job="kubernetes-service-endpoints", k8s_app="hubble", namespace="kube-system", node="kind-worker3", qtypes="A", query="qxmjmhckfuzcbozxjliglzigegc2vjcmv0igtletogu0vdukvus0vzmtizndu2.example.com.", service="hubble-metrics"}	43
hubble_dns_queries_total{app_kubernetes_io_managed_by="Helm", app_kubernetes_io_name="hubble", app_kubernetes_io_part_of="cilium", instance="172.18.0.3:9965", ips_returned="0", job="kubernetes-service-endpoints", k8s_app="hubble", namespace="kube-system", node="kind-worker3", qtypes="A", query="vghpcybpcybz21lihnlnnpdgl2zsbkyxrhroyxqgbmvlzhmgdg8gymugzxh.example.com.", service="hubble-metrics"}	43

Les metrics

...mais
suffisante
pour voir
une
anomalie



Les network flow logs

Ces logs offrent davantage d'informations

```
May 28 20:17:37.567: default/ubuntu-pod:55552 (ID:12263) -> 137.194.211.184:53 (world) policy-verdict:L4-Only EGRESS ALLOWED (UDP)
May 28 20:17:37.567: default/ubuntu-pod:55552 (ID:12263) -> 137.194.211.184:53 (world) to-proxy FORWARDED (UDP)
May 28 20:17:37.567: default/ubuntu-pod:55552 (ID:12263) -> 137.194.211.184:53 (world) dns-request proxy FORWARDED (DNS Query data.very.sensible.example.com. A)
May 28 20:17:37.567: default/ubuntu-pod:55552 (ID:12263) <- 137.194.211.184:53 (world) dns-response proxy FORWARDED (DNS Answer RCode: Non-Existent Domain TTL: 4294967295 (Proxy data.very.sensible.example.com. A))
May 28 20:17:37.568: default/ubuntu-pod:55552 (ID:12263) <- 137.194.211.184:53 (world) to-endpoint FORWARDED (UDP)
May 28 20:17:37.597: default/ubuntu-pod:35581 (ID:12263) -> 137.194.211.184:53 (world) policy-verdict:L4-Only EGRESS ALLOWED (UDP)
May 28 20:17:37.597: default/ubuntu-pod:35581 (ID:12263) -> 137.194.211.184:53 (world) to-proxy FORWARDED (UDP)
May 28 20:17:37.597: default/ubuntu-pod:35581 (ID:12263) -> 137.194.211.184:53 (world) dns-request proxy FORWARDED (DNS Query data.very.sensible.example.com. A)
May 28 20:17:37.598: default/ubuntu-pod:35581 (ID:12263) <- 137.194.211.184:53 (world) dns-response proxy FORWARDED (DNS Answer RCode: Non-Existent Domain TTL: 4294967295 (Proxy data.very.sensible.example.com. A))
May 28 20:17:37.598: default/ubuntu-pod:35581 (ID:12263) <- 137.194.211.184:53 (world) to-endpoint FORWARDED (UDP)
May 28 20:17:37.621: default/ubuntu-pod:57542 (ID:12263) -> 137.194.211.184:53 (world) policy-verdict:L4-Only EGRESS ALLOWED (UDP)
May 28 20:17:37.621: default/ubuntu-pod:57542 (ID:12263) -> 137.194.211.184:53 (world) to-proxy FORWARDED (UDP)
May 28 20:17:37.623: default/ubuntu-pod:57542 (ID:12263) -> 137.194.211.184:53 (world) dns-request proxy FORWARDED (DNS Query data.very.sensible.example.com. A)
May 28 20:17:37.624: default/ubuntu-pod:57542 (ID:12263) <- 137.194.211.184:53 (world) dns-response proxy FORWARDED (DNS Answer RCode: Non-Existent Domain TTL: 4294967295 (Proxy data.very.sensible.example.com. A))
May 28 20:17:37.624: default/ubuntu-pod:57542 (ID:12263) <- 137.194.211.184:53 (world) to-endpoint FORWARDED (UDP)
May 28 20:17:37.654: default/ubuntu-pod:51566 (ID:12263) -> 137.194.211.184:53 (world) policy-verdict:L4-Only EGRESS ALLOWED (UDP)
May 28 20:17:37.654: default/ubuntu-pod:51566 (ID:12263) -> 137.194.211.184:53 (world) to-proxy FORWARDED (UDP)
May 28 20:17:37.655: default/ubuntu-pod:51566 (ID:12263) -> 137.194.211.184:53 (world) dns-request proxy FORWARDED (DNS Query data.very.sensible.example.com. A)
May 28 20:17:37.655: default/ubuntu-pod:51566 (ID:12263) <- 137.194.211.184:53 (world) dns-response proxy FORWARDED (DNS Answer RCode: Non-Existent Domain TTL: 4294967295 (Proxy data.very.sensible.example.com. A))
May 28 20:17:37.655: default/ubuntu-pod:51566 (ID:12263) <- 137.194.211.184:53 (world) to-endpoint FORWARDED (UDP)
ubuntu@kind-2:~/Cilium$
```

Configuration d'Hubble exporter

Très simple à configurer !

Autres options :

- file rotation
- size limits,
- filters
- And field masks

```
! cilium-values.yaml
1  USER-SUPPLIED VALUES:
2  ∨ hubble:
3    enabled: true
4  ∨ metrics:
5    enableOpenMetrics: true
6    enabled:
7      - drop
8      - 'dns:query;sourceContext;identity;destinationContext:d
9      - tcp
10     - flow
11     - port-distribution
12     - icmp
13     - httpv2:expamplars=true
14  ∨ relay:
15    enabled: true
16  ∨ ui:
17    enabled: true
18  ∨ export:
19    ∨ static:
20      enabled: true
21      filePath: /var/run/cilium/hubble/events.log
22  ∨ operator:
23    ∨ prometheus:
24      enabled: true
25  ∨ prometheus:
26    enabled: true
27
```

Configuration d'Hubble exporter

- Attention : il y a un fichier de logs par nœud => ne pas se tromper de nœud !

```
ubuntu@kind-2:~$ hubble observe --protocol dns --since=30s
Jun 12 11:02:52.514: default/python-script-runner:40228 (ID:12263) -> 8.8.8.8:53 (world) dns-request proxy FORWARDED (DNS Query vghpcybpcybz21lihnlnbnpdgl2zs
bkyxrhihroyxqgbmvlzhmgdg8gymugzxh.example.com. A)
Jun 12 11:02:52.603: default/python-script-runner:40228 (ID:12263) <- 8.8.8.8:53 (world) dns-response proxy FORWARDED (DNS Answer RCode: Non-Existent Domain T
TL: 4294967295 (Proxy vghpcybpcybz21lihnlnbnpdgl2zsbkyxrhihroyxqgbmvlzhmgdg8gymugzxh.example.com. A))
Jun 12 11:02:52.605: default/python-script-runner:51693 (ID:12263) -> 8.8.8.8:53 (world) dns-request proxy FORWARDED (DNS Query mawx0cmf0zwqgb3zlcibetlmuckhlc
mugaxmgysbwyxnzd29yzdoguebzc3cwcm.example.com. A)
Jun 12 11:02:52.693: default/python-script-runner:51693 (ID:12263) <- 8.8.8.8:53 (world) dns-response proxy FORWARDED (DNS Answer RCode: Non-Existent Domain T
TL: 4294967295 (Proxy mawx0cmf0zwqgb3zlcibetlmuckhlc mugaxmgysbwyxnzd29yzdoguebzc3cwcm.example.com. A))
Jun 12 11:02:52.696: default/python-script-runner:45304 (ID:12263) -> 8.8.8.8:53 (world) dns-request proxy FORWARDED (DNS Query qxmjmhckfuzcbozxjliglzigegc2vj
cmv0igtletogu0vdukvus0vzmtizndu2.example.com. A)
Jun 12 11:02:52.783: default/python-script-runner:45304 (ID:12263) <- 8.8.8.8:53 (world) dns-response proxy FORWARDED (DNS Answer RCode: Non-Existent Domain T
TL: 4294967295 (Proxy qxmjmhckfuzcbozxjliglzigegc2vjcmv0igtletogu0vdukvus0vzmtizndu2.example.com. A))
ubuntu@kind-2:~$ kubectl -n kube-system exec ds/cilium -- tail -f /var/run/cilium/hubble/events.log | grep dns
Defaulted container "cilium-agent" out of: cilium-agent, config (init), mount-cgroup (init), apply-sysctl-overwrites (init), mount-bpf-fs (init), clean-cilium
-state (init), install-cni-binaries (init)
```

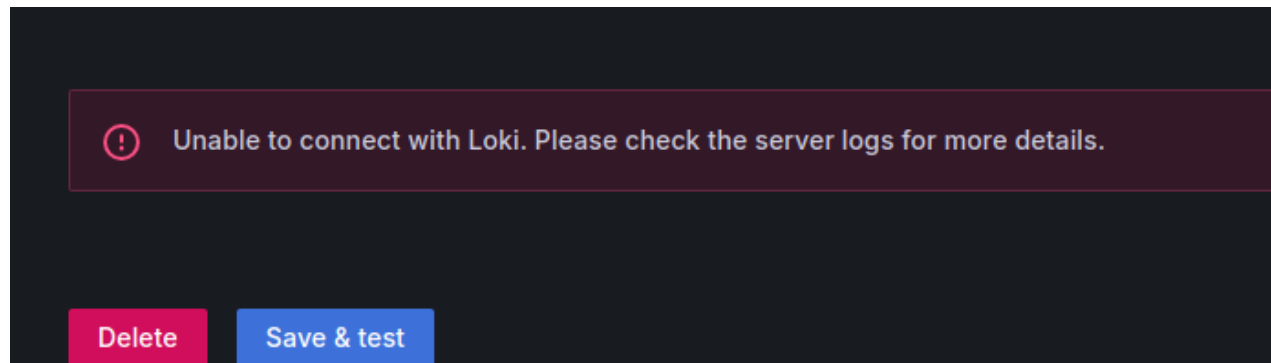
Ici, les logs DNS sont bien affichés par Hubble CLI, mais ils ne sont pas dans les logs que l'on consulte

Configuration de Loki

Utiliser un agrégateur de logs ? Loki !

```
helm upgrade --install --values loki.yaml loki grafana/loki-stack -n  
grafana-loki --create-namespace
```

Intégration dans grafana comme data source, mais...



Configuration de Loki

La réponse :

[Loki-stack] Bundled loki version is extremely out of date! #2873

New issue



MegaShinySnivy opened this issue on Dec 29, 2023 · 11 comments · Fixed by #2875



MegaShinySnivy commented on Dec 29, 2023

As per the title, the default loki version is extremely out of date, and cannot connect to modern grafana versions. It should be bumped to a more recent version, like 2.9.3. This bit me in [this issue \(#79846\)](#).



6

Assignees

No one assigned

Labels

None yet

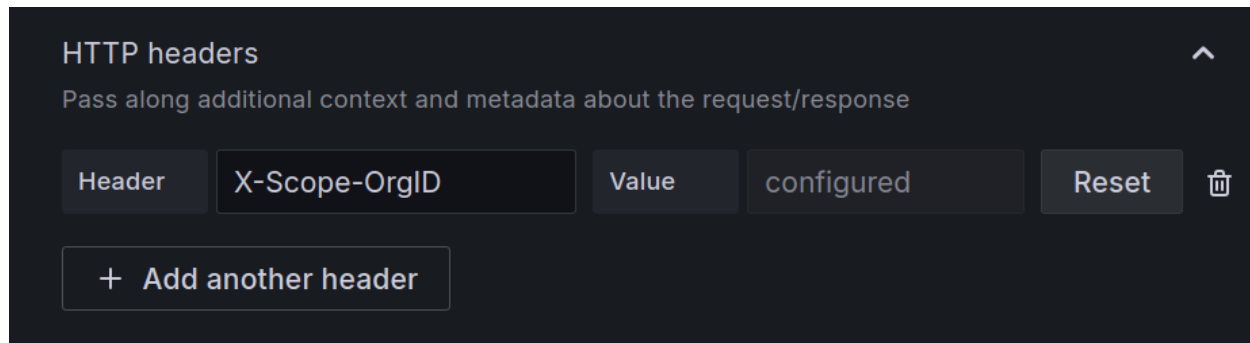
Projects

Configuration de Loki

Version la plus récente de Loki, compatible avec Grafana :

```
ubuntu@kind-2:~$ helm install --values Cilium/loki.yaml loki --namespace=monitoring grafana/loki
NAME: loki
LAST DEPLOYED: Tue Jun 11 17:35:42 2024
NAMESPACE: monitoring
STATUS: deployed
REVISION: 1
NOTES:
*****
Welcome to Grafana Loki
Chart version: 6.6.3
Chart Name: loki
Loki version: 3.0.0
*****
```

Ne pas oublier le HTTP header en plus dans Grafana:



The screenshot shows the 'HTTP headers' configuration panel in Grafana. At the top, it says 'HTTP headers' with an upward arrow icon, followed by the description 'Pass along additional context and metadata about the request/response'. Below this is a table with two columns: 'Header' and 'Value'. The first row contains 'X-Scope-OrgID' and 'configured'. To the right of the table are 'Reset' and 'trash' icons. At the bottom, there is a button that says '+ Add another header'.

Header	Value
X-Scope-OrgID	configured

+ Add another header

Configuration de Loki

Interface de configuration et de visualisation de Loki.

Label filters: namespace = kube-system, pod = cilium-qxd62, node_name = kind-worker3

Line contains: Text to find

Query: {namespace="kube-system", pod="cilium-qxd62", node_name="kind-worker3"} |= ''

Options: Type: Range, Line limit: 1000

Logs volume: Graph showing log volume over time (12:20 to 13:15). Legend: debug (blue), error (red). Total: 14 debug, 24 error.

Logs: Display results (Newest first, Oldest first). Common labels: cilium-agent, kube-system/cilium-agent, kind-worker3, cilium-qxd62. Line limit: 1000 (38 returned). Total bytes processed: 10.2 kB.

Log entries:

- > 2024-06-12 13:17:51.771 time="2024-06-12T11:17:51Z" level=error msg="Envoy: Version check failed" error="envoy version 801c612e442298b3be55f1a1089c44386570880d/1.29.4/Distribution/RELEASE/BoringSSL does not match with required version 8fccf45a8ab9da13824e0f14122d5db35673f3bb" subsys=envoy-proxy
- > 2024-06-12 13:17:51.771 time="2024-06-12T11:17:51Z" level=error msg="Timer job errored" module=agent.controlplane.envoy-proxy name=version-check func="envoy.registerEnvoyVersionCheck.func2 (pkg/envoy/cell.go:262)" error="envoy version 801c612e442298b3be55f1a1089c44386570880d/1.29.4/Distribution/RELEASE/BoringSSL does not match with required version 8fccf45a8ab9da13824e0f14122d5db35673f3bb"
- > 2024-06-12 13:17:51.771 time="2024-06-12T11:17:51Z" level=info msg="Envoy: Version 801c612e442298b3be55f1a1089c44386570880d/1.29.4/Distribution/RELEASE/BoringSSL" subsys=envoy-manager
- > 2024-06-12 13:12:51.770 time="2024-06-12T11:12:51Z" level=error msg="Timer job errored" module=agent.controlplane.envoy-proxy name=version-check func="envoy.registerEnvoyVersionCheck.func2 (pkg/envoy/cell.go:262)" error="envoy version 801c612e442298b3be55f1a1089c44386570880d/1.29.4/Distribution/RELEASE/BoringSSL does not match with required version 8fccf45a8ab9da13824e0f14122d5db35673f3bb"
- > 2024-06-12 13:12:51.770 time="2024-06-12T11:12:51Z" level=error msg="Envoy: Version check failed" error="envoy version 801c612e442298b3be55f1a1089c44386570880d/1.29.4/Distribution/RELEASE/BoringSSL does not match with required version 8fccf45a8ab9da13824e0f14122d5db35673f3bb"

Problème rencontré avec Promtail

- Promtail récupère les logs et les envoie à Loki, qui les agrège
=> un pod promtail par nœud
- On ajoute les logs d'Hubble en target

```
16  scrape_configs:
88
89      # New job for Hubble logs
90  - job_name: hubble-logs
91      kubernetes_sd_configs:          # we aim for pods
92      | - role: pod
93      |   relabel_configs:
94      |   | - source_labels:          # with the label k8s-app=cilium
95      |   | | - __meta_kubernetes_pod_label_k8s_app
96      |   |   regex: cilium
97      |   |   action: keep
98      |   | - source_labels:          # in the namespace kube-system
99      |   | | - __meta_kubernetes_namespace
100      |   |   regex: kube-system
101      |   |   action: keep
102      |   | - source_labels:          # we change the labels for readability
103      |   | | - __meta_kubernetes_namespace
104      |   |   action: replace
105      |   |   target_label: namespace
106      |   | - source_labels:
107      |   | | - __meta_kubernetes_pod_name
108      |   |   action: replace
109      |   |   target_label: pod
110      |   | - action: replace          # path of the logs
111      |   |   replacement: /var/run/cilium/hubble/events.log
112      |   |   target_label: __path__
113      |   pipeline_stages:
114      |   | - json:
115      |   |   expressions:
116      |   |   | qtype: flow.l7.dns.qtypes[0]
117      |   |   | query: flow.l7.dns.query
118      |   |   | source: flow.IP.source
119      |   |   | destination: flow.IP.destination
120      |   |   | direction: flow.traffic_direction
```

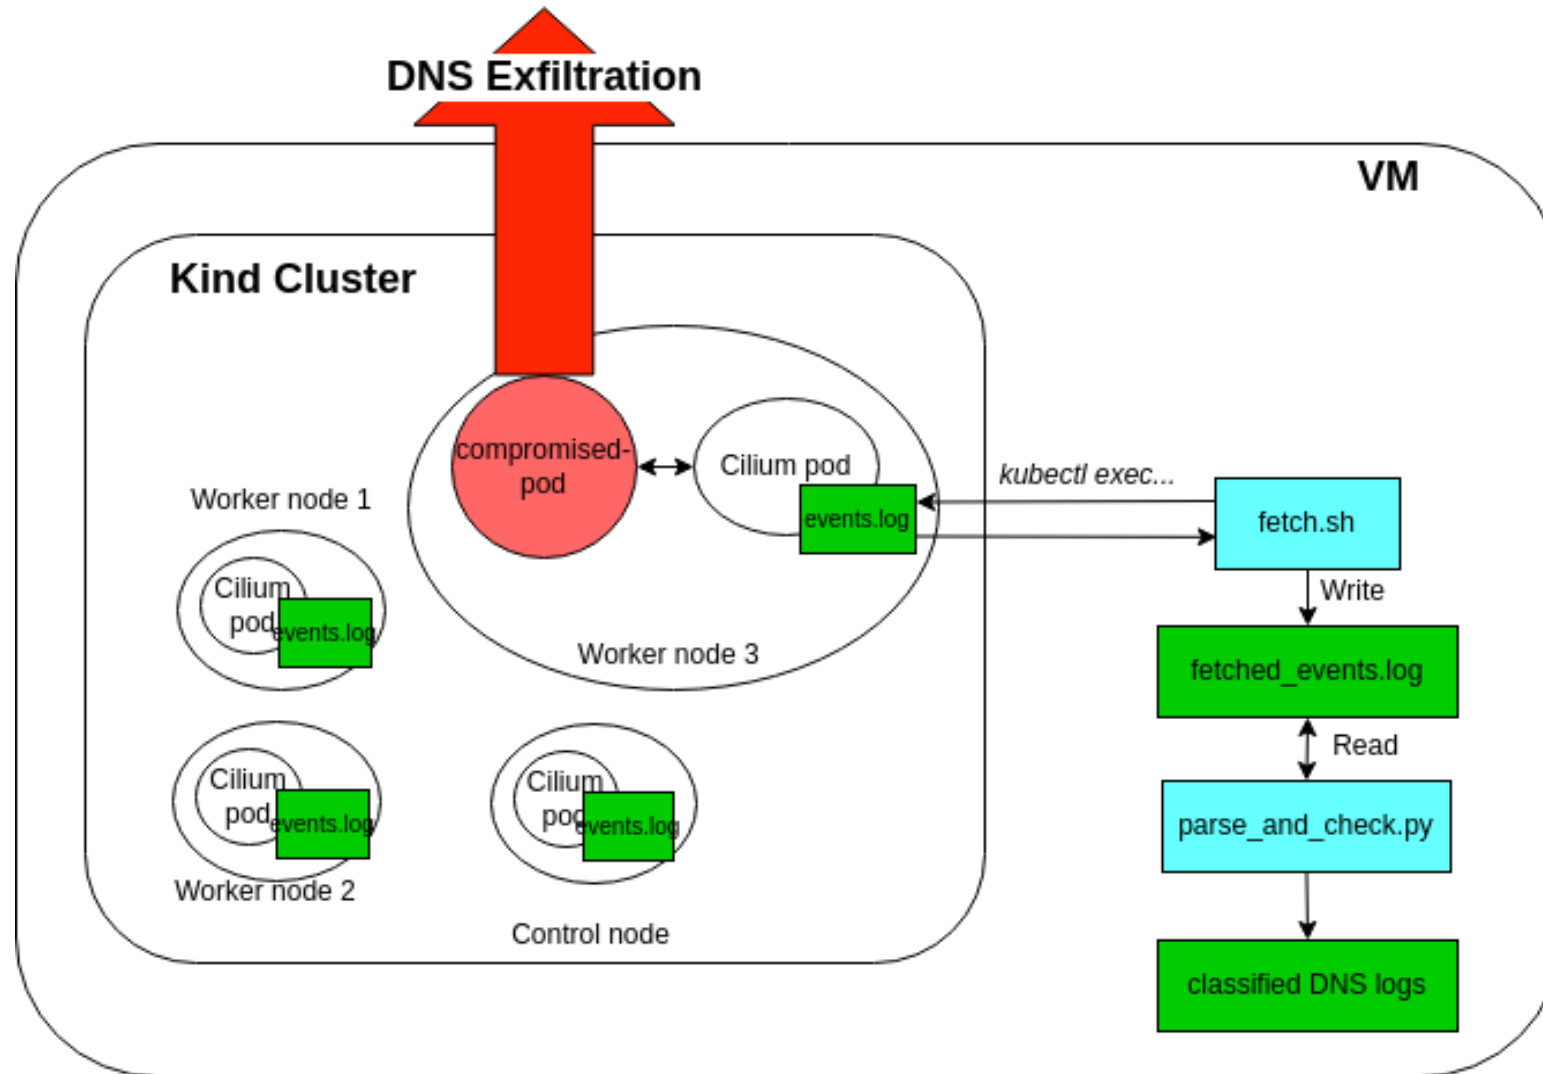
Problème rencontré avec Promtail

La target est bien ajoutée ...

```
level=info ts=2024-06-12T09:08:22.353308756Z caller=tailer.go:147 component=tailer msg="tail routine: started" path=/var/log/pods/kube-system_cilium-qxd62_ac2c173b-7c70-4557-8e22-8e8dc93bdb16/mount-cgroup/0.log
level=info ts=2024-06-12T09:08:22.353503983Z caller=filetarget.go:313 msg="watching new directory" directory=/var/log/pods/kube-system_cilium-qxd62_ac2c173b-7c70-4557-8e22-8e8dc93bdb16/mount-bpf-fs
level=info ts=2024-06-12T09:08:22.353775478Z caller=filetargetmanager.go:372 msg="Adding target" key="/var/run/cilium/hubble/events" log:{namespace="kube-system", pod="cilium-qxd62"}
level=info ts=2024-06-12T09:08:22.353780397Z caller=filetarget.go:313 msg="watching new directory" directory=/var/log/pods/kube-system_cilium-qxd62_ac2c173b-7c70-4557-8e22-8e8dc93bdb16/clean-cilium-state
level=info ts=2024-06-12T09:08:22.353933634Z caller=filetarget.go:313 msg="watching new directory" directory=/var/log/pods/kube-system_cilium-qxd62_ac2c173b-7c70-4557-8e22-8e8dc93bdb16/install-cni-binaries
```

Cependant, les logs ne remontent pas à Loki

Récupération des logs pour classification



Récupération des logs pour classification

Démo !

DNS exfiltration classification : Splunk's approach



Products ▾

Solutions ▾

Why Splunk? ▾

Resources ▾

Company ▾

Support ▾



Free Splunk

Splunk Blogs

Security

DevOps

Artificial Intelligence

Platform

Leadership

Partners

.conf

Splunk Life

More ▾

Security

JULY 07, 2023 | 8 MINUTE READ

Machine Learning in Security: Detect DNS Data Exfiltration Using Deep Learning



DNS exfiltration classification : Splunk's approach

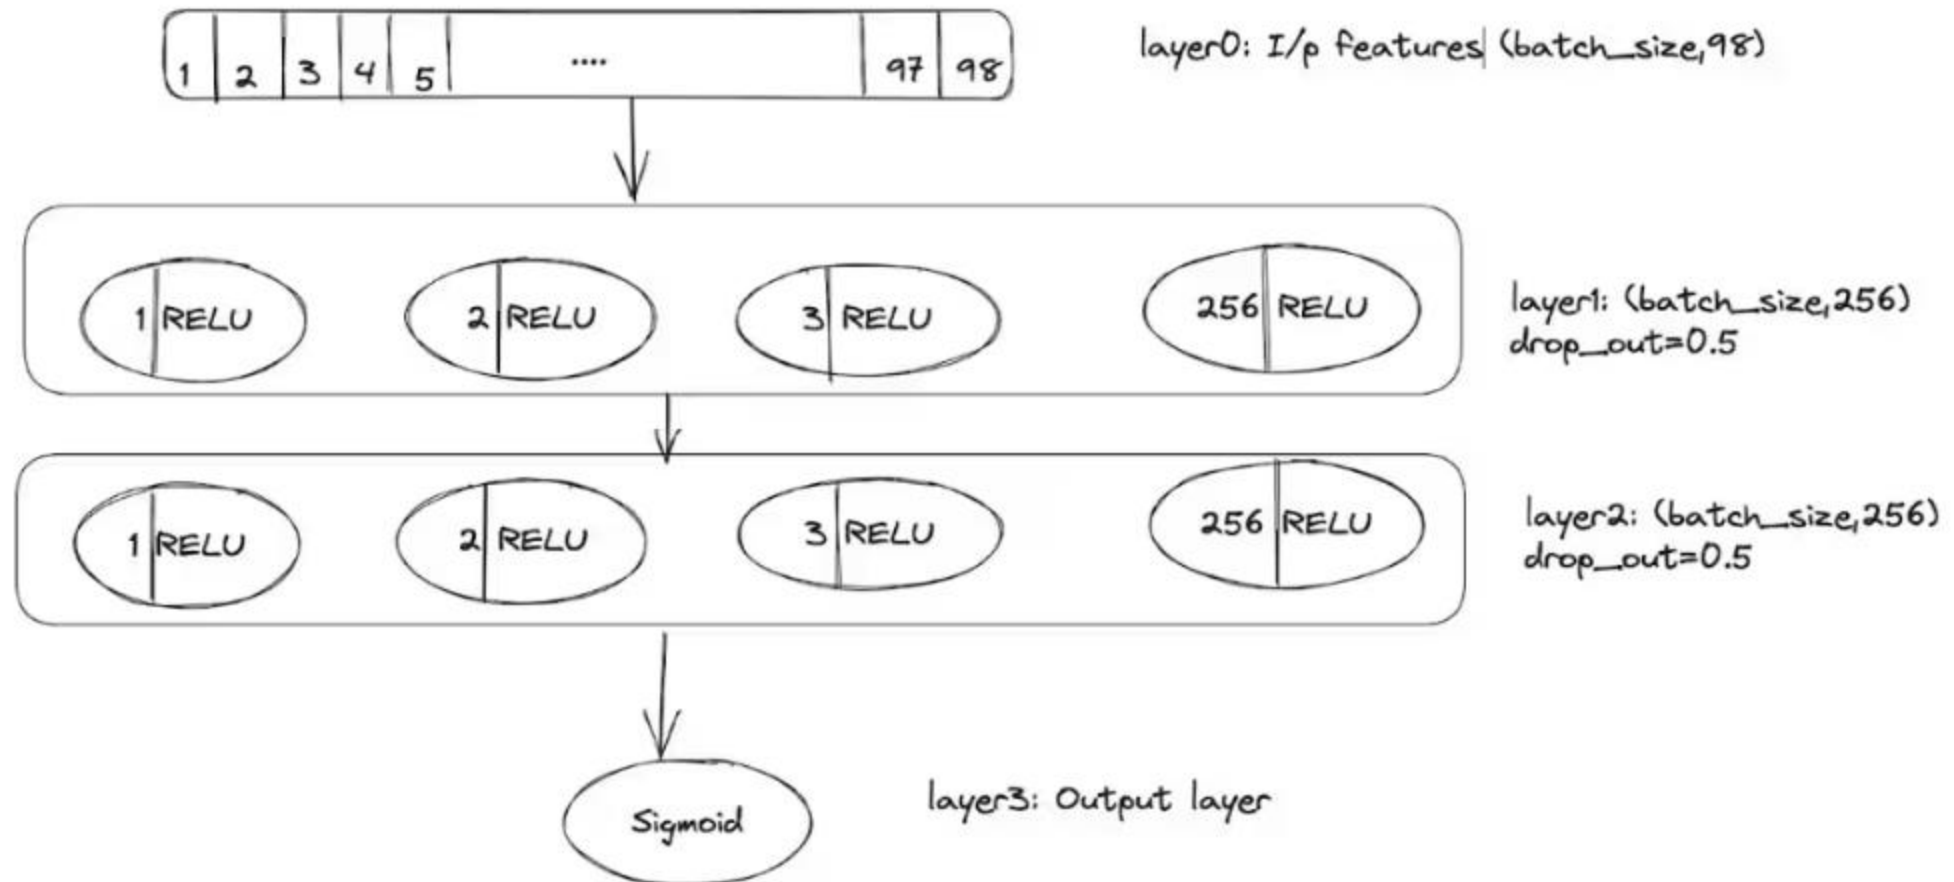
High volume DNS exfiltrations are easy to spot (e.g. with our Grafana dashboard !)

But low volume DNS exfiltrations require to look at DNS queries one by one => the following elements are taken into account :


- Tokenized DNS request (94 characters)
- Length
- Entropy
- Average length over sliding window (past events between (source, dns_server))
- Average entropy over sliding window (past events between (source, dns_server))

['0', '1', '2', '3', '4', '5', '6', '7', '8', '9', 'a', 'b', 'c', 'd', 'e', 'f', 'g', 'h', 'i', 'j', 'k', 'l', 'm', 'n', 'o', 'p', 'q', 'r', 's', 't', 'u', 'v', 'w', 'x', 'y', 'z', 'A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J', 'K', 'L', 'M', 'N', 'O', 'P', 'Q', 'R', 'S', 'T', 'U', 'V', 'W', 'X', 'Y', 'Z', '!', '"', '#', '\$', '%', '&', "'", '(', ')', '*', '+', ',', '-', '.', '/', ':', ';', '<', '=', '>', '?', '@', '[', '\\', ']', '^', '_', '`', '{', '|', '}', '~']

DNS exfiltration classification : Splunk's approach



DNS exfiltration classification : Splunk's approach



Splunk, 2023, Confusion Matrix

Un dataset trouvé



DNS Exfiltration Dataset

Published: 11 July 2023 | Version 3 | DOI: 10.17632/c4n7fckkz3.3

Contributors: Kristijan Ziza, Pavle Vuletić, Predrag Tadić

Description

DNS exfiltration dataset was recorded in a realistic network environment. More than 50 million DNS requests were recorded on one of the ISP's DNS servers. The data in the dataset was anonymised by changing all IP addresses using injective mapping.

Features in the dataset are split into single request and aggregate features. Single request or DNS label-based features can be calculated for each DNS request independently using only the textual characteristics of the request. On the other hand, aggregate features are calculated using multiple subsequent request from one client to a particular TLD. This reduces the size of the dataset to about 35 million records. The complete list of features with descriptions can be found in dataset_description.txt file. For all of the features which are based on finding English words in the request we used about 60.000 most common English words. The list of used words can be found in english_words.txt.

The main dataset (dataset.csv) contains regular requests and exfiltrations performed using DNSExfiltrator and Iodine tools. Additional dataset (dataset_modified.csv) contains only exfiltrations executed with modified DNSExfiltrator tool. Waiting times between two consecutive requests in this dataset are randomised and the requests also have lower entropy causing the detection to be much harder.

Perspective ML

- Inspecter le dataset
- Tester un NN simple sur le dataset, et voir les performances

To do list

1. Entrainement ML
2. Intégration dans "parse and classify"
3. Hubble et les traces ? Démo trouvée, à tester
4. Revenir sur les limites des metrics (nombre de labels max, pk que certains, etc.)
5. Cilium policy rule pour bloquer certaines exfiltrations